



▼ Importing Libraries

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn import metrics
from sklearn import svm
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn import preprocessing
from sklearn.metrics import confusion_matrix, classification_report
from mpl_toolkits import mplot3d
from sklearn.multiclass import OneVsOneClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import BernoulliNB
```

▼ Reading & Analyzing Data

```
df = pd.read_csv("/content/drug200.csv")

df.head()
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	DrugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	DrugY

```
df.describe()
```

	Age	Na_to_K
count	200.000000	200.000000
mean	44.315000	16.084485
std	16.544315	7.223956
min	15.000000	6.269000
25%	31.000000	10.445500
50%	45.000000	13.936500
75%	58.000000	19.380000
max	74.000000	38.247000

```
df.isnull().sum()
```

```
Age          0
Sex          0
BP           0
Cholesterol  0
Na_to_K      0
Drug         0
dtype: int64
```

```
df.isna().sum()
```

```
Age          0
Sex          0
BP           0
Cholesterol  0
Na_to_K      0
Drug         0
dtype: int64
```

- No missing value
- 6 columns
- 200 rows

Variable Description

- Age: Age of patient
- Sex: Gender of patient
- BP: Blood pressure of patient
- Cholesterol: Cholesterol of patient
- Na_to_K: Sodium to Potassium Ratio in Blood
- Drug: Drug Type

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Age         200 non-null   int64
1   Sex         200 non-null   object
2   BP          200 non-null   object
3   Cholesterol 200 non-null   object
4   Na_to_K     200 non-null   float64
5   Drug        200 non-null   object
dtypes: float64(1), int64(1), object(4)
memory usage: 9.5+ KB
```

- float64(1): Na_to_K
- int64(1): Age
- object(4): Sex, BP, Cholesterol, Drug

Univariate Variable Analysis


Age Variable

```
print("Max Age:", df.Age.max())
print("Min Age:", df.Age.min())
```

```
Max Age: 74
Min Age: 15
```

```
# Age distribution
plt.figure(figsize = (9,5))
sns.distplot(df.Age)
plt.show()
```

/usr/local/lib/python3.8/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Use `displot` instead.



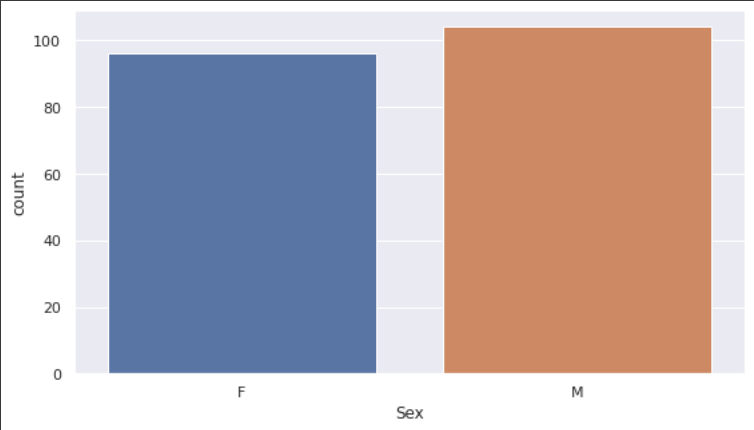
- Age range is between 15 and 74.

Sex Variable

```
df.Sex.value_counts()

M      104
F       96
Name: Sex, dtype: int64

# Sex Distribution
plt.figure(figsize=(9,5))
sns.countplot(x = df.Sex)
plt.show()
```




- The ratio of gender seems balanced in the data
- This is a categorical variable. It would be better if we apply label encoder to avoid any error during model implementation.

BP Variable

```
df.BP.value_counts()

HIGH      77
LOW       64
NORMAL    59
Name: BP, dtype: int64

plt.figure(figsize = (9,5))
sns.countplot(df.BP)
plt.show()
```




Cholesterol Variable

```
df.Cholesterol.value_counts()

HIGH      103
NORMAL    97
Name: Cholesterol, dtype: int64

plt.figure(figsize = (9,5))
sns.countplot(df.Cholesterol)
plt.show()
```

```
/usr/local/lib/python3.8/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From warnings.warn(
```



- Cholesterol is a balanced data.
- It is categorical and label encoder will apply on it.



▼ Na_to_K Variable




```
print("Max Na_to_K:",df.Na_to_K.max())
print("Min Na_to_K:",df.Na_to_K.min())
print("Mean Na_to_K:",df.Na_to_K.mean())
```

```
Max Na_to_K: 38.247
Min Na_to_K: 6.269
Mean Na_to_K: 16.084485
```

```
plt.figure(figsize = (9,5))
sns.distplot(df.Na_to_K)
plt.show()
```

```
/usr/local/lib/python3.8/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be re
warnings.warn(msg, FutureWarning)
```




▼ Drug Variable

```
df.Drug.value_counts()
```

```
DrugY      91
drugX      54
drugA      23
drugC      16
drugB      16
Name: Drug, dtype: int64
```

```
plt.figure(figsize = (9,5))
sns.countplot(df.Drug)
plt.show()
```

```
/usr/local/lib/python3.8/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From warnings.warn(
```



- Drug is target column and you can see that it is unbalanced dataset. Using K Fold cross-validation would be better for reliable results.


▼ Basic Data Analysis

- Age – Drug
- Sex – Drug

- BP -- Drug
- Cholesterol -- Drug

Age -- Drug

```
plt.figure(figsize = (9,5))
sns.swarmplot(x = "Drug", y = "Age",data = df)
plt.legend(df.Drug.value_counts().index)
plt.title("Age -- Drug")
plt.show()
```



```
print("Minimum Age of DrugB",df.Age[df.Drug == "drugB"].min())
print("Maximum Age of DrugA",df.Age[df.Drug == "drugA"].max())
```

Minimum Age of DrugB 51
Maximum Age of DrugA 50


- DrugB is taken only by older than 51 years old.
- DrugA is taken only by younger than 50 years old.

Sex -- Drug

```
df_Sex_Drug = df.groupby(["Drug","Sex"]).size().reset_index(name = "Count")
df_Sex_Drug
```

	Drug	Sex	Count
0	DrugY	F	47
1	DrugY	M	44
2	drugA	F	9
3	drugA	M	14
4	drugB	F	6
5	drugB	M	10
6	drugC	F	7
7	drugC	M	9
8	drugX	F	27
9	drugX	M	27

```
plt.figure(figsize = (9,5))
sns.barplot(x = "Drug",y="Count", hue = "Sex",data = df_Sex_Drug)
plt.title("Sex -- Drug")
plt.show()
```



- Male people get drugA, drugB and drugC more than male people.
- Female people get DrugY more than female people.
- drugX seems equal for male and female people.


- According to this graph, Sex feature is not an important feature for classification.

BP -- Drug

```
df_BP_Drug = df.groupby(["Drug","BP"]).size().reset_index(name = "Count")
df_BP_Drug
```

	Drug	BP	Count
0	DrugY	HIGH	38
1	DrugY	LOW	30
2	DrugY	NORMAL	23
3	drugA	HIGH	23
4	drugB	HIGH	16
5	drugC	LOW	16
6	drugX	LOW	18
7	drugX	NORMAL	36

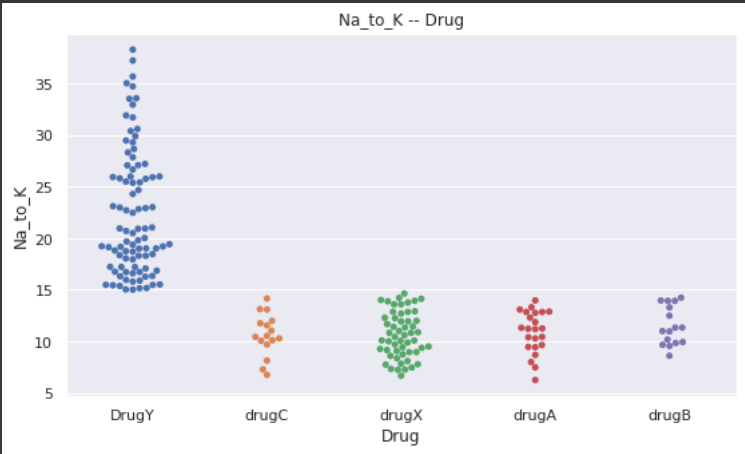
```
plt.figure(figsize = (9,5))
sns.barplot(x = "Drug",y="Count", hue = "BP",data = df_BP_Drug)
plt.title("BP -- Drug")
plt.show()
```



- drugA and drugB are got only by people who have HIGH blood pressure.
- drugC is got by people who have LOW blood pressure.
- drugX is got by people who have HIGH blood pressure.
- BP is an important feature for classification.

Na_to_K -- Drug

```
plt.figure(figsize = (9,5))
sns.swarmplot(x = "Drug", y = "Na_to_K",data = df)
plt.title("Na_to_K -- Drug")
plt.show()
```



```
print("Minimum Na_to_K for DrugY:",df.Na_to_K[df.Drug == "DrugY"].min())
```

Minimum Na_to_K for DrugY: 15.015

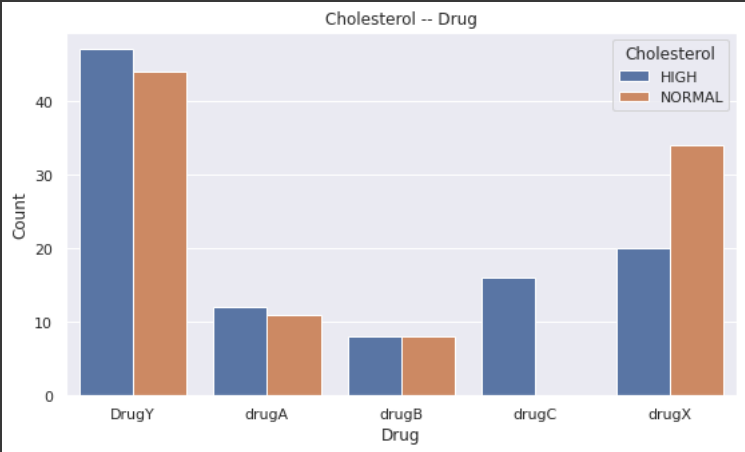
- People who have Na_to_K ratio is bigger than 15, get DrugY.
- We can create a new feature from here.

Cholesterol -- Drug

```
df_CH_Drug = df.groupby(["Drug","Cholesterol"]).size().reset_index(name = "Count")
df_CH_Drug
```

	Drug	Cholesterol	Count
0	DrugY	HIGH	47
1	DrugY	NORMAL	44
2	drugA	HIGH	12
3	drugA	NORMAL	11
4	drugB	HIGH	8
5	drugB	NORMAL	8
6	drugC	HIGH	16
7	drugX	HIGH	20
8	drugX	NORMAL	34


```
plt.figure(figsize = (9,5))
sns.barplot(x = "Drug",y="Count", hue = "Cholesterol",data = df_CH_Drug)
plt.title("Cholesterol -- Drug")
plt.show()
```



- drugC is got by people who have HIGH cholesterol.
- Cholesterol is an important feature to classify drugC

Na_to_K -- BP -- Drug

```
plt.figure(figsize = (9,5))
sns.swarmplot(x = "Drug", y = "Na_to_K",hue="BP",data = df)
plt.legend()
plt.title("Na_to_K -- BP -- Drug")
plt.show()
```



- If people have HIGH blood pressure and Na_to_K ratio is lower than 15 , they get drugA and drugB only.
- If people have LOW blood pressure and Na_to_K ratio is lower than 15 , they get drugC only.

Preparing Data and Feature Engineering

Create New Features

Na_to_K_Bigger_Than_15

If Na_to_K is bigger than 15, it is always drugY.

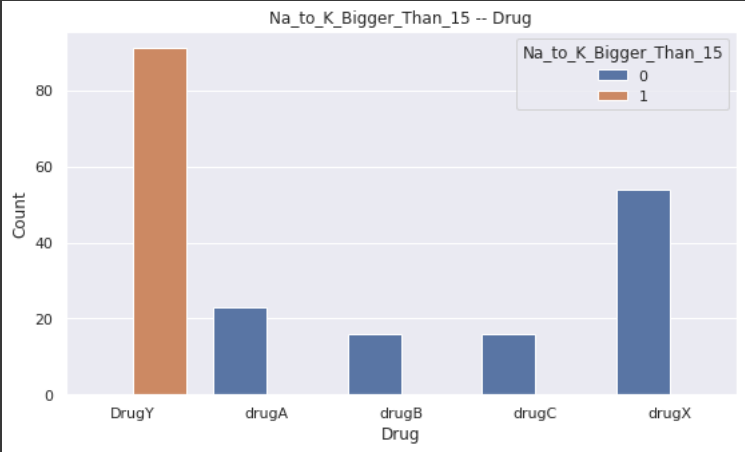
```
df['Na_to_K_Bigger_Than_15'] = [1 if i >=15.015 else 0 for i in df.Na_to_K]
df.head()
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug	Na_to_K_Bigger_Than_15
0	23	F	HIGH	HIGH	25.355	DrugY	1
1	47	M	LOW	HIGH	13.093	drugC	0
2	47	M	LOW	HIGH	10.114	drugC	0
3	28	F	NORMAL	HIGH	7.798	drugX	0
4	61	F	LOW	HIGH	18.043	DrugY	1

```
df_NaK15 = df.groupby(["Drug","Na_to_K_Bigger_Than_15"]).size().reset_index(name = "Count")
df_NaK15
```

	Drug	Na_to_K_Bigger_Than_15	Count
0	DrugY	1	91
1	drugA	0	23
2	drugB	0	16
3	drugC	0	16
4	drugX	0	54

```
plt.figure(figsize = (9,5))
sns.barplot(x = "Drug",y="Count", hue = "Na_to_K_Bigger_Than_15",data = df_NaK15)
plt.title("Na_to_K_Bigger_Than_15 -- Drug")
plt.show()
```



- Na_to_K_Bigger_Than_15 feature will be important feature to drugY classification.

Label Encoding

We will convert from object to int64

- Sex
- BP
- Cholesterol
- Na_to_K
- Na_to_K_Bigger_Than_15

```
from sklearn.preprocessing import LabelEncoder

def label_encoder(y):
    le = LabelEncoder()
    df[y] = le.fit_transform(df[y])
```

```
label_list = ["Sex","BP","Cholesterol","Na_to_K","Na_to_K_Bigger_Than_15","Drug"]

for l in label_list:
    label_encoder(l)
```

```
df.head()
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug	Na_to_K_Bigger_Than_15
0	23	0	0	0	167	0	1
1	47	1	1	0	89	3	0
2	47	1	1	0	43	3	0
3	28	0	2	0	10	4	0
4	61	0	1	0	133	0	1

Train Test Split

```
from sklearn.model_selection import train_test_split

x = df.drop(["Drug"],axis=1)
```



```
y = df.Drug

x_train, x_test, y_train, y_test = train_test_split(x,y,test_size = 0.2, random_state = 42, shuffle = True)

y_train = y_train.values.reshape(-1,1)
y_test = y_test.values.reshape(-1,1)

print("x_train shape:",x_train.shape)
print("x_test shape:",x_test.shape)
print("y_train shape:",y_train.shape)
print("y_test shape:",y_test.shape)
x_train shape: (160, 6)
x_test shape: (40, 6)
y_train shape: (160, 1)
y_test shape: (40, 1)
```

Data was splitted as 80% train data and 20% test data.

▼ KNN Classifier

To find best score of KNN model, We will try different value of n_neighbors, p, and weights parameters.

▼ Default Parameters

```
from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier()
knn.fit(x_train,y_train)
knnpred=knn.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test, knnpred))

Accuracy: 0.65
/usr/local/lib/python3.8/dist-packages/sklearn/neighbors/_classification.py:198: DataConversionWarning: A column-vector y was passed wh
return self._fit(X, y)

print (metrics.classification_report(y_test, knnpred))
```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	15
1	0.38	0.50	0.43	6
2	0.50	0.67	0.57	3
3	0.50	0.20	0.29	5
4	0.50	0.45	0.48	11
accuracy			0.65	40
macro avg	0.56	0.56	0.55	40
weighted avg	0.65	0.65	0.64	40

▼ Random Forest

To find best score of Random Forest model, we will try different value of n_estimators and criterion parameters.

▼ Default Parameters

```
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier(random_state = 42)
rfc.fit(x_train,y_train)
rf_pred = rfc.predict(x_test)
print("Accuracy:",metrics.accuracy_score(y_test, rf_pred))

Accuracy: 0.975
<ipython-input-40-f188aa6b6aaf>:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the
rfc.fit(x_train,y_train)

print (metrics.classification_report(y_test, rf_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	15
1	1.00	1.00	1.00	6
2	1.00	1.00	1.00	3
3	1.00	0.80	0.89	5
4	0.92	1.00	0.96	11
accuracy			0.97	40
macro avg	0.98	0.96	0.97	40
weighted avg	0.98	0.97	0.97	40

▼ SVM Classifier

To find best score of SVM model, we will try different value of C, kernel, degree and gamma parameters. The easy way to do this is GridSearchCV method.

▼ Default Parameters

```
from sklearn.svm import SVC
svc = SVC(random_state = 42)
svc.fit(x_train,y_train)

prediction=svc.predict(x_test)
print('The accuracy of the SVM is:',metrics.accuracy_score(prediction,y_test))

The accuracy of the SVM is: 0.65
/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array
y = column_or_1d(y, warn=True)
```

```
print (metrics.classification_report(y_test, prediction))
```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	15
1	0.00	0.00	0.00	6
2	0.00	0.00	0.00	3
3	0.00	0.00	0.00	5
4	0.46	1.00	0.63	11
accuracy			0.65	40
macro avg	0.28	0.40	0.32	40
weighted avg	0.48	0.65	0.54	40

```
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-d
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-d
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-d
_warn_prf(average, modifier, msg_start, len(result))
```

▼ Logistic Regression

```
model = LogisticRegression()
model.fit(x_train, y_train)
y_predicted = model.predict(x_test)
print('The accuracy of the Logistic Regression is',metrics.accuracy_score(y_predicted,y_test))

The accuracy of the Logistic Regression is 0.85
/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array
y = column_or_1d(y, warn=True)
/usr/local/lib/python3.8/dist-packages/sklearn/linear_model/_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
n_iter_i = _check_optimize_result(
```

```
print(classification_report(y_test,y_predicted ))
```

	precision	recall	f1-score	support
0	0.88	1.00	0.94	15
1	1.00	0.83	0.91	6
2	1.00	1.00	1.00	3
3	0.00	0.00	0.00	5
4	0.73	1.00	0.85	11
accuracy			0.85	40
macro avg	0.72	0.77	0.74	40
weighted avg	0.76	0.85	0.80	40

```
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-d
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-d
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-d
_warn_prf(average, modifier, msg_start, len(result))
```

▼ Naive Bayes

```
Nb= BernoulliNB()
Nb.fit(x_train, y_train)
y_predict_nb = Nb.predict(x_test)
nbs=metrics.accuracy_score(y_predict_nb,y_test)
print('The accuracy of the Naive Bayes is',metrics.accuracy_score(y_predict_nb,y_test))

The accuracy of the Naive Bayes is 0.8
/usr/local/lib/python3.8/dist-packages/sklearn/utils/validation.py:993: DataConversionWarning: A column-vector y was passed when a 1d array
y = column_or_1d(y, warn=True)
```

```
print(classification_report(y_test,y_predict_nb ))
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

```
0      1.00      1.00      1.00      15
1      0.67      1.00      0.80      6
2      0.00      0.00      0.00      3
3      0.00      0.00      0.00      5
4      0.69      1.00      0.81      11

accuracy              0.80      40
macro avg            0.47      0.60      0.52      40
weighted avg         0.66      0.80      0.72      40

/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-d
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-d
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-d
_warn_prf(average, modifier, msg_start, len(result))
```

Conclusion

```
print('Naive Bayes:      ',nbs)
print('KNN:              ',metrics.accuracy_score(y_test, knnpred))
print('Logistic Regression: ',metrics.accuracy_score(y_predicted,y_test))
print('SVM:              ',metrics.accuracy_score(prediction,y_test))
print('Random Forest:    ',metrics.accuracy_score(y_test, rf_pred))
```

Naive Bayes: 0.8
KNN: 0.65
Logistic Regression: 0.85
SVM: 0.65
Random Forest: 0.975